

Introduction

January 17, 2025

```
[ ]: '''  
    Reinforcement Learning -->  
  
    Reinforcement Learning (RL) is a type of machine learning where  
    an agent learns how to behave in an environment by performing actions  
    and receiving rewards or penalties. The goal of the agent is to learn  
    a policy that maximizes the cumulative reward over time.  
  
    Key Components of Reinforcement Learning -->  
  
    Agent: The learner or decision-maker.  
    Environment: Everything the agent interacts with.  
    State (S): A representation of the current situation of the agent in the  
    ↪environment.  
    Action (A): A set of all possible moves the agent can make.  
    Reward (R): Feedback from the environment to evaluate an action.  
    Policy ( ): A strategy that the agent uses to determine actions based on  
    ↪the current state.  
    Value Function (V(s)): Predicts the long-term reward of a state under a  
    ↪policy.  
    Q-Function (Q(s,a)): Predicts the long-term reward of taking action a in  
    ↪state  
    s and following the policy thereafter.  
    '''
```

```
[ ]: '''  
    Workflow of Reinforcement Learning -->  
  
    The agent observes the current state of the environment.  
    Based on the policy, it takes an action.  
    The environment transitions to a new state and provides a reward.  
    The agent updates its policy based on the reward and the observed  
    ↪transition.  
    This loop continues until the agent learns an optimal policy.  
    '''
```

[]: '''

In the context of Reinforcement Learning, Upper Confidence Bound (UCB) and Thompson Sampling (TS) are commonly used for tackling the multi-armed bandit (MAB) problem. These methods aim to balance exploration (trying out

↳ less-tested

actions) and exploitation (choosing actions that seem to perform best based on current information).

Multi-Armed Bandit Problem -->

You are given multiple "arms" (choices/actions), each with an unknown probability distribution of rewards.

The goal is to pull the arms in such a way that maximizes the cumulative reward over time.

Upper Confidence Bound (UCB) -->

UCB is a deterministic approach to address the exploration-exploitation

↳ trade-off.

It assumes that the true mean reward of an arm is within a certain

↳ confidence

interval based on the number of times the arm has been pulled.

Workflow:

At each time step, calculate the UCB for all arms.

Select the arm with the highest UCB.

Update the rewards and the pull count for the selected arm.

Advantages:

Guarantees logarithmic regret.

Effective when the distribution of rewards has clear separation.

Limitations:

Works best with stationary reward distributions.

Computationally expensive for a large number of arms due to repeated

↳ calculations.

Thompson Sampling -->

Thompson Sampling is a probabilistic approach based on Bayesian inference.

It assumes a prior distribution for each arm's reward probabilities and updates this posterior based on observed rewards.

Workflow:

Initialize a prior distribution for each arm

(e.g., Beta distribution for binary rewards).

For each arm, sample from its posterior distribution.

*Select the arm with the highest sampled value.
Observe the reward and update the posterior distribution for the selected_*
↪ arm.

Advantages:

*Naturally balances exploration and exploitation through random sampling.
Handles non-stationary rewards better than UCB.*

Limitations:

*Computationally more intensive for complex reward distributions.
Requires well-defined priors for reward distributions.*

'''

[]: '''

Applications -->

Both methods are widely used in :

*Recommendation Systems (e.g., which ad to show a user).
Clinical Trials (e.g., testing multiple treatments).
Online Platforms (e.g., optimizing website layouts).*

'''

Feature	Upper Confidence Bound	Thompson Sampling
Exploration Strategy	Selects action with highest UCB	Samples from posterior
Exploitation	Balances exploration and exploitation	Uses Bayesian inference
Computation	Deterministic	Stochastic
Adaptability	Slower to adapt	Quickly adapts with more data
Assumptions	Requires known reward range	Handles uncertainty better
Complexity	Moderate	Can be computationally intensive
Applications	Good for known environments	Useful in uncertain environments