

Upper-Confidence-Bound

January 17, 2025

```
[ ]: '''  
    Upper Confidence Bound (UCB) -->  
  
    The Upper Confidence Bound (UCB) is a popular approach in multi-armed  
    bandit problems where an agent must balance exploration (trying different  
    options to gather information) and exploitation (choosing the best-known_  
    ↪option).  
    It is used to optimize the selection of actions based on observed rewards.  
    '''
```

```
[ ]: '''  
    Problem Context -->  
  
    Imagine you have  $n$  slot machines (or "arms") with unknown probabilities  
    of reward. The goal is to maximize your total reward over time by  
    pulling the arms intelligently.  
  
    How It Works -->  
  
    Initialization:  
    Pull each arm once to get initial reward estimates.  
  
    Iterate:  
    For each arm, compute the UCB score.  
    Select the arm with the highest UCB score.  
    Update the reward estimate and pull count for the chosen arm.  
    Repeat until a stopping condition (e.g., time limit) is met.  
  
    Key Features -->  
  
    Exploration:  
    Arms with fewer pulls are favored because of the  $\ln t/N_i(t)$  term.  
  
    Exploitation:  
    Arms with higher average rewards ( $\mu_i$ ) are preferred over time.  
    '''
```

```
[ ]: '''
    Applications -->

    Recommender Systems: Selecting items or advertisements to display.
    Clinical Trials: Choosing treatments to test on patients.
    Online Learning: Optimizing decisions in real-time scenarios.
    '''
```

```
[3]: # Importing Libraries -->

import pandas as pd
import numpy as np
import math
import matplotlib.pyplot as plt
```

```
[4]: # Importing Dataset -->

data = pd.read_csv('Data/Ads_CTR_Optimisation.csv')
data.head(10)
```

```
[4]:
```

	Ad 1	Ad 2	Ad 3	Ad 4	Ad 5	Ad 6	Ad 7	Ad 8	Ad 9	Ad 10
0	1	0	0	0	1	0	0	0	1	0
1	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	0
3	0	1	0	0	0	0	0	1	0	0
4	0	0	0	0	0	0	0	0	0	0
5	1	1	0	0	0	0	0	0	0	0
6	0	0	0	1	0	0	0	0	0	0
7	1	1	0	0	1	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	1	0	0	0	0	0	0	0

```
[5]: # Implementing UCB -->

num_users = 10000
num_ads = 10
ads_selected = []
num_selections = [0] * num_ads
sum_reward = [0] * num_ads
total_reward = 0

for rounds in range(0,num_users):
    ad = 0
    max_upper_bound = 0

    for ads in range(0,num_ads):
```

```

    if (num_selections[ads] > 0):
        avg_reward = sum_reward[ads] / num_selections[ads]
        delta = math.sqrt(3/2 * math.log(rounds+1)/num_selections[ads])
        upper_bound = avg_reward + delta

    else:
        upper_bound = 1e400 # Super high value

    if (upper_bound > max_upper_bound):
        max_upper_bound = upper_bound
        ad = ads

ads_selected.append(ad)
num_selections[ad] += 1
reward = data.values[rounds, ad]
sum_reward[ad] += reward
total_reward += reward


```

```

[6]: # Visualizing The Results -->

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of time ad selected')
plt.show()

```



```
[ ]: # Ad 5 is selected the most !
```

```
[7]: # Let's see how quickly it can identify -->

# Changing num_users to 1000 to reduce number of rounds !

num_users = 1000
num_ads = 10
ads_selected = []
num_selections = [0] * num_ads
sum_reward = [0] * num_ads
total_reward = 0

for rounds in range(0,num_users):
    ad = 0
    max_upper_bound = 0

    for ads in range(0,num_ads):

        if (num_selections[ads] > 0):
```

```

        avg_reward = sum_reward[ads] / num_selections[ads]
        delta = math.sqrt(3/2 * math.log(rounds+1)/num_selections[ads])
        upper_bound = avg_reward + delta

    else:
        upper_bound = 1e400 # Super high value

    if (upper_bound > max_upper_bound):
        max_upper_bound = upper_bound
        ad = ads

    ads_selected.append(ad)
    num_selections[ad] += 1
    reward = data.values[rounds, ad]
    sum_reward[ad] += reward
    total_reward += reward


```

```
[8]: # Visualizing The Results -->
```

```

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of time ad selected')
plt.show()

```



```
[ ]: # It is still identifying the ad which is most selected !
```

```
[9]: # Let's make it 500 -->
```

```
num_users = 500
num_ads = 10
ads_selected = []
num_selections = [0] * num_ads
sum_reward = [0] * num_ads
total_reward = 0

for rounds in range(0, num_users):
    ad = 0
    max_upper_bound = 0

    for ads in range(0, num_ads):

        if (num_selections[ads] > 0):
            avg_reward = sum_reward[ads] / num_selections[ads]
            delta = math.sqrt(3/2 * math.log(rounds+1)/num_selections[ads])
```

```

        upper_bound = avg_reward + delta

    else:
        upper_bound = 1e400 # Super high value

    if (upper_bound > max_upper_bound):
        max_upper_bound = upper_bound
        ad = ads

    ads_selected.append(ad)
    num_selections[ad] += 1
    reward = data.values[rounds, ad]
    sum_reward[ad] += reward
    total_reward += reward


```

[10]: # Visualizing The Results -->

```

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of time ad selected')
plt.show()

```



[]: # *You can see it is unable to identify the ad in 500 rounds !*