

Thompson-Sampling

January 17, 2025

[]: '''

Thompson Sampling -->

Thompson Sampling is a popular reinforcement learning and decision-making algorithm used for solving the multi-armed bandit problem. It is designed to balance the trade-off between exploration (trying new options to gain more information) and exploitation (choosing the option that currently seems best).

'''

[]: '''

How Thompson Sampling Works -->

Thompson Sampling uses a probabilistic approach to determine which arm to pull at each step. It models the uncertainty of rewards using a prior distribution and updates this distribution based on observed outcomes (Bayesian updating).

Assume Prior Distributions :

Initially, assume a prior distribution (e.g., Beta distribution for Bernoulli rewards) for each arm's probability of reward.

Sampling :

For each arm, sample a value from its posterior distribution.

The posterior combines prior beliefs with the data observed so far.

Select Arm:

Choose the arm with the highest sampled value (argmax of samples).

Observe Reward:

Pull the chosen arm, observe the reward, and update the posterior distribution for that arm.

Repeat:

Continue this process, balancing exploration and exploitation automatically as the posterior distributions converge.

'''

```
[ ]: '''
    Key Advantages -->

    Efficient Balance: Thompson Sampling inherently balances
    exploration and exploitation.
    Bayesian Approach: Uses probabilistic reasoning to model uncertainty.
    Adaptable: Works with various reward distributions.

    Applications -->

    Online Advertising: Optimizing click-through rates by choosing
    which ad to display.
    Recommendation Systems: Deciding which products to recommend to users.
    Clinical Trials: Testing new drugs or treatments.
    '''
```

```
[4]: # Importing Libraries -->

import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
```

```
[5]: # Importing Dataset -->

data = pd.read_csv('Data/Ads_CTR_Optimisation.csv')
data.head(10)
```

```
[5]:
```

	Ad 1	Ad 2	Ad 3	Ad 4	Ad 5	Ad 6	Ad 7	Ad 8	Ad 9	Ad 10
0	1	0	0	0	1	0	0	0	1	0
1	0	0	0	0	0	0	0	0	1	0
2	0	0	0	0	0	0	0	0	0	0
3	0	1	0	0	0	0	0	1	0	0
4	0	0	0	0	0	0	0	0	0	0
5	1	1	0	0	0	0	0	0	0	0
6	0	0	0	1	0	0	0	0	0	0
7	1	1	0	0	1	0	0	0	0	0
8	0	0	0	0	0	0	0	0	0	0
9	0	0	1	0	0	0	0	0	0	0

```
[8]: # Implementing Thompson Sampling -->

num_users = 10000
num_ads = 10
ads_selected = []
num_rewards_1 = [0] * num_ads
num_rewards_0 = [0] * num_ads
```

```

total_reward = 0

for rounds in range(0, num_users):
    ad = 0
    max_random = 0

    for ads in range(0, num_ads):
        random_beta = random.betavariate(num_rewards_1[ads] + 1,
        num_rewards_0[ads] + 1)

        if (random_beta > max_random):
            max_random = random_beta
            ad = ads

    ads_selected.append(ad)
    reward = data.values[rounds, ad]

    if (reward == 1):
        num_rewards_1[ad] += 1

    else:
        num_rewards_0[ad] += 1

    total_reward += reward

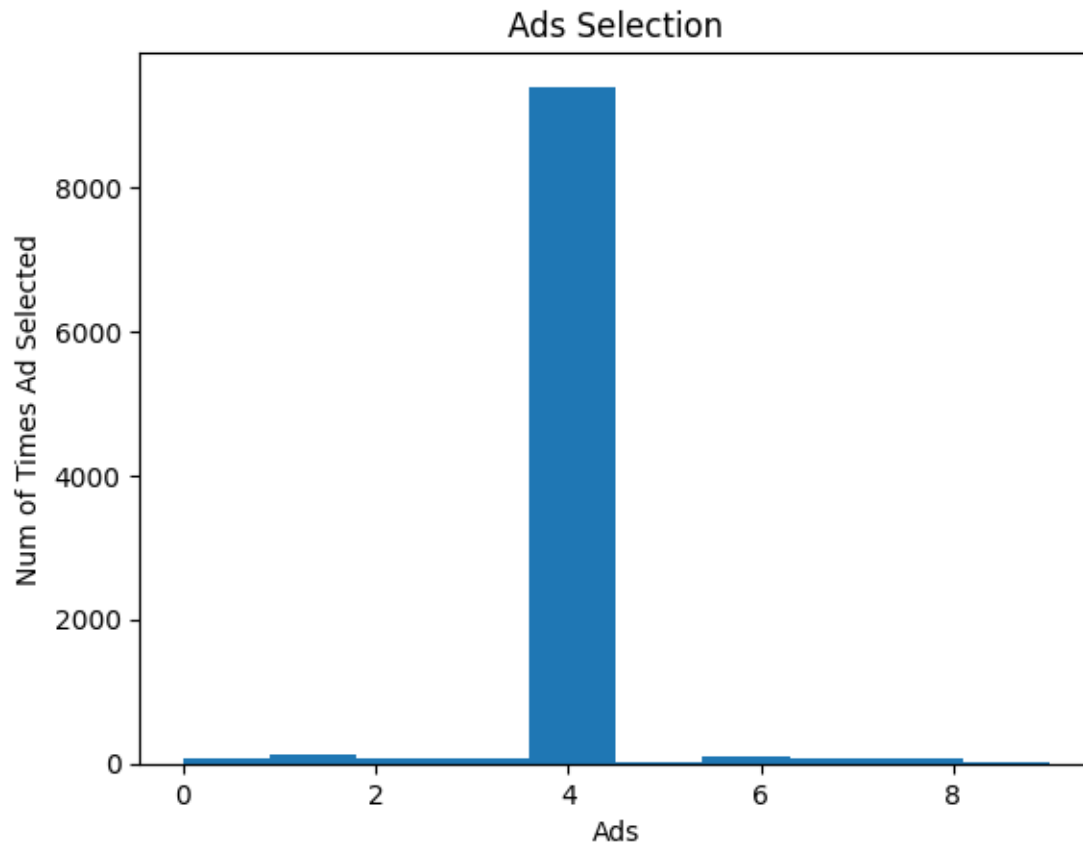
```

```

[9]: # Visualzing The Results -->

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of Times Ad Selected')
plt.show()

```



[10]: *# Let's Try It With 1000 Rounds -->*

```
num_users = 1000
num_ads = 10
ads_selected = []
num_rewards_1 = [0] * num_ads
num_rewards_0 = [0] * num_ads
total_reward = 0

for rounds in range(0, num_users):
    ad = 0
    max_random = 0

    for ads in range(0, num_ads):
        random_beta = random.betavariate(num_rewards_1[ads] + 1,
        ↪ num_rewards_0[ads] + 1)

        if (random_beta > max_random):
            max_random = random_beta
```

```

        ad = ads

ads_selected.append(ad)
reward = data.values[rounds, ad]

if (reward == 1):
    num_rewards_1[ad] += 1

else:
    num_rewards_0[ad] += 1

total_reward += reward


```

[11]: # Visualzing The Results -->

```

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of Times Ad Selected')
plt.show()

```



```
[15]: #   Let's Try It With 500 Rounds -->

num_users = 200
num_ads = 10
ads_selected = []
num_rewards_1 = [0] * num_ads
num_rewards_0 = [0] * num_ads
total_reward = 0

for rounds in range(0, num_users):
    ad = 0
    max_random = 0

    for ads in range(0, num_ads):
        random_beta = random.betavariate(num_rewards_1[ads] + 1,
↪num_rewards_0[ads] + 1)

        if (random_beta > max_random):
            max_random = random_beta
            ad = ads

    ads_selected.append(ad)
    reward = data.values[rounds, ad]

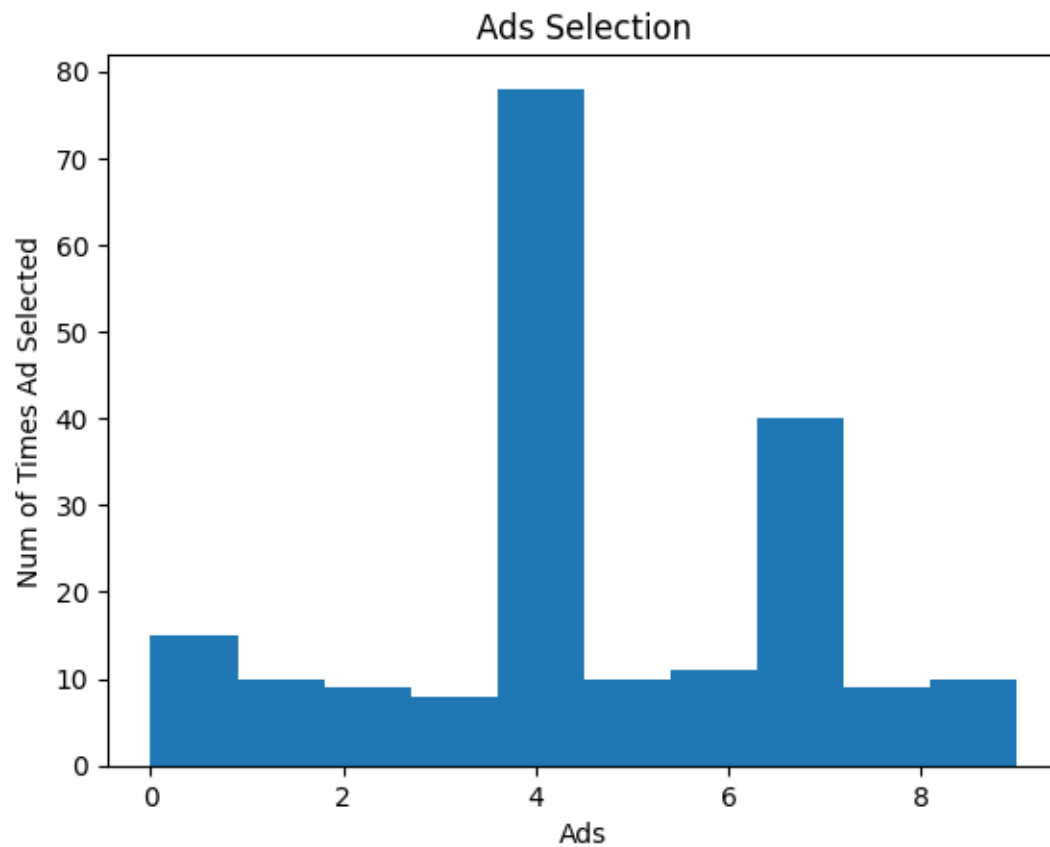
    if (reward == 1):
        num_rewards_1[ad] += 1

    else:
        num_rewards_0[ad] += 1

    total_reward += reward
```

```
[16]: #   Visualzing The Results -->

plt.hist(ads_selected)
plt.title('Ads Selection')
plt.xlabel('Ads')
plt.ylabel('Num of Times Ad Selected')
plt.show()
```



```
[ ]: '''  
    Thompson Sampling is Efficient as it is able to identify  
    the most selected ad in just around 200 rounds while UCB  
    was unable to identify in around 500 rounds.  
    '''
```