

CNN

February 20, 2025

[]: '''

Convolutional Neural Networks -->

Introduction -->

Convolutional Neural Networks (CNNs) are specialized deep learning models designed to process grid-like data, particularly images. They mimic the way the human brain processes visual information, recognizing patterns, edges, textures, and hierarchical structures automatically.

'''

[]: '''

Intuition Behind CNNs -->

Why Not Traditional Neural Networks ?

Traditional fully connected neural networks (FCNNs) struggle with high-dimensional image data due to:

Parameter Explosion: Fully connected layers require enormous numbers of weights when handling high-resolution images.

Spatial Redundancy: Pixels in an image have local dependencies;

FCNNs fail to leverage these relationships efficiently.

Translation Variance: A simple shift in an image can affect FCNN predictions drastically.

CNNs address these issues by:

Using local receptive fields instead of fully connected layers.

Sharing weights across the image using kernels.

Applying downsampling (pooling) to reduce dimensions while preserving essential features.

'''

[]: '''

Core Components of CNNs -->

1. Convolutional Layer

The core idea of CNNs lies in convolution operations, which apply a small filter (kernel) over the input image to extract essential features such as edges, corners, and textures.

Filters (Kernels): Small matrices that slide over the image.

Stride: Defines the step size for moving the filter.

Padding: Controls how borders of the image are handled (zero-padding or valid-padding).

Feature Maps: The output of convolution layers after applying activation functions.

2. Activation Function (ReLU)

ReLU introduces non-linearity, allowing CNNs to learn complex patterns.

$\text{ReLU}(x) = \max(0, x)$

3. Pooling Layer

Pooling reduces spatial dimensions while retaining significant information.

Max Pooling: Retains the highest value in a region.

Average Pooling: Takes the average of pixel values in a region.

Pooling helps with:

Reducing computation.

Enhancing feature robustness.

Mitigating overfitting.

4. Fully Connected Layer

After extracting meaningful features, CNNs flatten the feature maps and pass them through dense layers to make predictions.

5. Softmax (Output Layer)

Converts logits into probability distributions for classification.

'''

[]: '''

Applications of CNNs -->

Image Classification (e.g., CIFAR-10, ImageNet)

Object Detection (e.g., YOLO, SSD)

```
Facial Recognition
Medical Image Analysis
Autonomous Vehicles
'''
```

```
[39]: # Importing Libraries -->

import numpy as np
import tensorflow as tf
from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Input, Conv2D, Dense, Flatten,
↳MaxPooling2D, Dropout

print("Libraries Imported !")
```

Libraries Imported !

```
[40]: # Image Augmentation [Training Set] -->

train_datagen = ImageDataGenerator(
    rescale = 1./255,
    shear_range = 0.2,
    zoom_range = 0.2,
    horizontal_flip = True
)

training_set = train_datagen.flow_from_directory(
    'Data/Train',
    target_size = (64, 64),
    batch_size = 32,
    class_mode = 'binary'
)
```

Found 8000 images belonging to 2 classes.

```
[41]: # Image Augmentation [Test Set] -->

test_datagen = ImageDataGenerator(
    rescale = 1./255
)

test_set = test_datagen.flow_from_directory(
    'Data/Test',
    target_size = (64, 64),
    batch_size = 32,
    class_mode = 'binary'
)
```

```
)
```

Found 2000 images belonging to 2 classes.

```
[42]: # Building The CNN -->

model = Sequential([
    Input(shape=(64,64,3)),
    Conv2D(filters=32, kernel_size=3, activation='relu'),
    MaxPooling2D(pool_size=2, strides=2),
    Flatten(),
    Dense(units=128, activation='relu'),
    Dense(units=1, activation='sigmoid')
])

print("Model Built !")
```

Model Built !

```
[43]: # Summarizing The Model -->

model.summary()
```

Model: "sequential_2"

Layer (type)	Output Shape	
Param #		
conv2d_2 (Conv2D)	(None, 62, 62, 32)	
↪ 896		
max_pooling2d_2 (MaxPooling2D)	(None, 31, 31, 32)	
↪ 0		
flatten_2 (Flatten)	(None, 30752)	
↪ 0		
dense_4 (Dense)	(None, 128)	
↪ 3,936,384		
dense_5 (Dense)	(None, 1)	
↪ 129		

Total params: 3,937,409 (15.02 MB)

Trainable params: 3,937,409 (15.02 MB)

Non-trainable params: 0 (0.00 B)

```
[44]: #   Compiling The Model -->

model.compile(optimizer='adam', loss='binary_crossentropy',
              metrics=['accuracy'])

print("Model Compiled !")
```

Model Compiled !

```
[46]: #   Training The Model -->

model.fit(
    x = training_set,
    validation_data = test_set,
    epochs = 25
)
```

Epoch 1/25

250/250 147s 588ms/step -

accuracy: 0.5650 - loss: 0.7286 - val_accuracy: 0.6850 - val_loss: 0.6020

Epoch 2/25

250/250 59s 234ms/step -

accuracy: 0.6760 - loss: 0.6050 - val_accuracy: 0.6840 - val_loss: 0.5908

Epoch 3/25

250/250 57s 227ms/step -

accuracy: 0.7007 - loss: 0.5746 - val_accuracy: 0.7210 - val_loss: 0.5473

Epoch 4/25

250/250 81s 225ms/step -

accuracy: 0.7228 - loss: 0.5492 - val_accuracy: 0.7090 - val_loss: 0.5640

Epoch 5/25

250/250 83s 227ms/step -

accuracy: 0.7326 - loss: 0.5304 - val_accuracy: 0.7125 - val_loss: 0.5655

Epoch 6/25

250/250 57s 227ms/step -

accuracy: 0.7414 - loss: 0.5197 - val_accuracy: 0.6980 - val_loss: 0.6099

Epoch 7/25

250/250 61s 245ms/step -

accuracy: 0.7482 - loss: 0.5013 - val_accuracy: 0.7430 - val_loss: 0.5394

Epoch 8/25

250/250 78s 228ms/step -

accuracy: 0.7658 - loss: 0.4814 - val_accuracy: 0.7605 - val_loss: 0.5189

Epoch 9/25

250/250 61s 244ms/step -

accuracy: 0.7706 - loss: 0.4734 - val_accuracy: 0.7365 - val_loss: 0.5420
 Epoch 10/25
 250/250 56s 225ms/step -
 accuracy: 0.7676 - loss: 0.4699 - val_accuracy: 0.7635 - val_loss: 0.5221
 Epoch 11/25
 250/250 82s 225ms/step -
 accuracy: 0.7865 - loss: 0.4480 - val_accuracy: 0.7610 - val_loss: 0.5144
 Epoch 12/25
 250/250 83s 230ms/step -
 accuracy: 0.7832 - loss: 0.4500 - val_accuracy: 0.7710 - val_loss: 0.5115
 Epoch 13/25
 250/250 82s 229ms/step -
 accuracy: 0.7959 - loss: 0.4357 - val_accuracy: 0.7720 - val_loss: 0.5070
 Epoch 14/25
 250/250 81s 226ms/step -
 accuracy: 0.7976 - loss: 0.4240 - val_accuracy: 0.7710 - val_loss: 0.5227
 Epoch 15/25
 250/250 81s 223ms/step -
 accuracy: 0.8122 - loss: 0.4058 - val_accuracy: 0.7750 - val_loss: 0.5124
 Epoch 16/25
 250/250 58s 230ms/step -
 accuracy: 0.8294 - loss: 0.3862 - val_accuracy: 0.7520 - val_loss: 0.5687
 Epoch 17/25
 250/250 59s 235ms/step -
 accuracy: 0.8313 - loss: 0.3797 - val_accuracy: 0.7500 - val_loss: 0.5987
 Epoch 18/25
 250/250 62s 248ms/step -
 accuracy: 0.8371 - loss: 0.3684 - val_accuracy: 0.7720 - val_loss: 0.5321
 Epoch 19/25
 250/250 61s 244ms/step -
 accuracy: 0.8525 - loss: 0.3290 - val_accuracy: 0.7545 - val_loss: 0.5417
 Epoch 20/25
 250/250 61s 243ms/step -
 accuracy: 0.8589 - loss: 0.3265 - val_accuracy: 0.7630 - val_loss: 0.5772
 Epoch 21/25
 250/250 57s 228ms/step -
 accuracy: 0.8538 - loss: 0.3291 - val_accuracy: 0.7745 - val_loss: 0.5574
 Epoch 22/25
 250/250 58s 233ms/step -
 accuracy: 0.8527 - loss: 0.3285 - val_accuracy: 0.7565 - val_loss: 0.6065
 Epoch 23/25
 250/250 80s 227ms/step -
 accuracy: 0.8745 - loss: 0.2939 - val_accuracy: 0.7560 - val_loss: 0.6208
 Epoch 24/25
 250/250 82s 228ms/step -
 accuracy: 0.8656 - loss: 0.3037 - val_accuracy: 0.7530 - val_loss: 0.6938
 Epoch 25/25
 250/250 57s 229ms/step -

accuracy: 0.8811 - loss: 0.2796 - val_accuracy: 0.7630 - val_loss: 0.6182

[46]: <keras.src.callbacks.history.History at 0x2c0ab61dc60>

[51]: # Accuracy -->

```
test_loss, test_acc = model.evaluate(test_set)
print(f"Test Accuracy: {test_acc * 100:.2f}%")
```

63/63 6s 99ms/step -
accuracy: 0.7739 - loss: 0.6049
Test Accuracy: 76.30%

[50]: # Predicting Results For Single Image -->

```
test_image = image.load_img('Data/Sample/cat_or_dog_1.jpg', target_size=(64,64))
test_image = image.img_to_array(test_image)
test_image = np.expand_dims(test_image, axis=0)

result = model.predict(test_image)
training_set.class_indices

if (result[0][0] == 1) :
    print("It is a Dog !")
else :
    print("It is a Cat !")
```

1/1 0s 140ms/step
It is a Dog !

[52]: # Predicting Results For Test Set -->

```
true_labels = test_set.classes

predictions = model.predict(test_set) # This returns probabilities
predicted_labels = np.argmax(predictions, axis=1) # Convert to class indices

correct_predictions = np.sum(predicted_labels == true_labels)
total_images = len(true_labels)

accuracy = (correct_predictions / total_images) * 100

print(f"Correct Predictions: {correct_predictions}/{total_images}")
print(f"\nTest Accuracy: {accuracy:.2f}%")
```

63/63 6s 89ms/step
Correct Predictions: 1000/2000

Test Accuracy: 50.00%